

Impact of Variations of Similarity and Prediction Techniques on User-based and Item-based Collaborative Filtering Recommendations

ANH HOANG, Davidson College, USA

MENGFAN WANG, Davidson College, USA

MIKE REMEZO, Davidson College, USA

MALAVIKA KALANI, Davidson College, USA

Given a wide array of implementations of recommender systems that were developed as a result of the exponential growth of the recommender system-related research field, it is of great importance to be able to measure these different algorithms against each other and perform an objective analysis of their functionality. In this paper, we dismantle the two most popular recommender methods, collaborative filtering and content-based methods, into separate components to evaluate each of them when accompanied by different similarity and prediction parameters. We seek to propose the best prediction algorithm - one that we define to be the recommending implementation that demonstrates the highest level of accuracy and coverage across all algorithms analyzed.

Additional Key Words and Phrases: Recommender Systems, Collaborative Filtering, Similarity and Prediction techniques

ACM Reference Format:

Anh Hoang, Mengfan Wang, Mike Remezo, and Malavika Kalani. 2022. Impact of Variations of Similarity and Prediction Techniques on User-based and Item-based Collaborative Filtering Recommendations. In *16th ACM Conference on Recommender Systems, Seattle, WA USA, Sept. 18 - 23, 2022*. ACM, New York, NY, USA, 15 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

1 INTRODUCTION

Owing to the booming growth of e-commerce in the last decade, there has been a drastic movement from the physical to virtual space with regards to consumers' habits. Digital platforms such as Netflix and Amazon to name a few, have become an integral part of every individual's life as these platforms guarantee leisure, convenience, and efficiency. They are able to give their users a more personalized experience because of a strong network of algorithms known as *Recommender Systems*, which solve the problem of how to filter, prioritize, and efficiently deliver suggestions on items that best fit consumers' needs.

There has been extensive research into the field of recommender systems with substantial focus placed on its two major paradigms: collaborative filtering and content-based methods. Collaborative filtering generates recommendations for a user through a systematic analysis of data collected from other users sharing similar tastes or information needs. Through the use of this recommender method, the recommending process can be conducted on certain types of items that may prove to be hard to analyze with computers such as feelings, ideas, and movies. In addition, collaborative filtering allows for serendipitous recommendations, which means that it may provide users with suggestions on items beyond their expectations. Collaborative filtering also improves the quality of its recommendations as more interactions are recorded over time. On the other hand, content-based methods examine items similar to other items that a user

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2022 Association for Computing Machinery.

Manuscript submitted to ACM

has liked in the past. This recommender method compensates for collaborative filtering's cold start problem, which means that the user-based algorithm requires a certain amount of data for it to work effectively; however, content-based methods don't show potential for increasing effectiveness over time since they can only operate on users' existing interest. Overall, both methods still present imperfections in their mechanisms; this particular problem has created an intellectual venue for a vast amount of research literature that looks into different ways of implementing a recommender system.

In light of the overwhelming number of existing recommender algorithms, this particular problem warrants critical analysis of these different implementations. The core recommending algorithm can be broken down into three steps for both user-based and item-based methods: (1) Apply different weightings to all users/ items to calculate the similarity among users and items (2) Create a subset of users/items to predict the ratings for an item that the user has not rated (3) Perform normalization of ratings and calculate the prediction using a weighted array of selected neighbors' ratings[1]. During each step in the process, we can apply different similarity and prediction parameters, such as Pearson correlation and Euclidean distance for the similarity weighting, or we can choose to apply significance weighting to account for only items or users that have a very high similarity score. Given that each implementation will output different results with varying accuracy and coverage, the question of what combinations of similarity and prediction techniques will produce the best results is worthy of discussion.

Our paper presents an algorithmic evaluation of content-based and user-based recommender systems through different implementations of the core algorithm, which allows for fine-grained control over distinct variations of similarity and prediction parameters. Ultimately, our goal is to propose the best recommendation algorithm based on two metrics: accuracy and coverage.

2 RELATED WORK

In this section, we briefly present some of the research literature related to the early implementations of collaborative filtering systems as well as works related to the functioning of user and item-based recommendations.

Recommender systems have become a popular system that most companies and platforms rely on for offering the best recommendations to users through content-based and collaborative filtering methods. Since its initial development, collaborative filtering methods have offered more advantages than content-based filtering methods as stated by [1]. Collaborative filtering filters items based on quality and user preferences and also provides serendipitous recommendations to users as suggested by [1] in his paper *An Algorithmic Framework for Performing Collaborative Filtering*.

GroupLens, a distributed system for using ratings from other users to predict an active user's ratings in articles, was the first to introduce an automated collaborative filtering system. GroupLens mainly used the ratings to determine which users' ratings were most similar to each other and it used ratings from similar users to predict how well a user would like a certain article. Better Bit Bureaus (one of the entities in GroupLens' netnews system) are servers that are used to predict ratings. In a user-based collaborative filtering system, we can calculate a predicted rating of an item not rated by a user by finding other users similar to that user who has rated that item. In a similar manner, *Better Bit Bureaus* uses the opinions of other people who have already rated the articles and combines those ratings to generate a predicted rating.

The tremendous advancement in the use of technology has led to an increase in the number of users on websites and digital platforms. The research paper [2] describes how item-based collaborative filtering systems can enable recommender systems to function at a larger scale by using the correlation between different items to make recommendations for the user. In our paper, we have shown the use of the Pearson correlation and the Euclidian distance method to

compute the similarity between two items. A few research studies have also focused on other methods to compute the similarity between two items. [2] describes two other methods: cosine-based similarity (considers two items as vectors and computes the similarity by finding the cosine of the angle between them) and adjusted cosine similarity (takes into account the difference in rating scale between different users by subtracting the respective user average from each co-rated pair). [1] lays a strong foundation for our research as it provides a similar analysis of different algorithms to identify the most accurate recommender system. Herlocker's detailed explanation of each component of the collaborative filtering process proves to be an integral driving force behind our findings.

3 THESIS

In this section, we will describe the research that we performed in this work and its main objective. We will also provide a brief description of how the research was conducted along with important details regarding the mathematical ideas and formulae used. Collaborative filtering is gradually becoming a reliable technique for managing the increasing load of user information while also ensuring a strong recommendation system for users to have a more personalized experience. Our research aims to apply different variations on user-based and item-based collaborative filtering processes. We will analyze the respective results to identify a set of recommender systems algorithm that produces the best recommendations to users. For each of the user-based and item-based methods, we will be producing results using the following similarity weighting measures:

- (1) *Euclidian distance* calculates the similarity between any two users across all the items they have rated in common.

In recommender systems, the *euclidian distance* is a value ranging between 0 and 1 where 0 indicates no similarity and 1 indicates high similarity. The formula for calculating similarity between the ratings of two users (u_1, u_2) for two items (i_1, i_2) using euclidian distance is shown below:

$$\text{Euclidian similarity distance} = \frac{1}{1 + \sqrt{(u_1.i_1 - u_2.i_1)^2 + (u_1.i_2 - u_2.i_2)^2}} \quad (1)$$

- (2) *Pearson Correlation* also calculates the similarity between two users across all the items they have rated in common. The *pearson correlation coefficient* is a value ranging between -1 and $+1$. A value of -1 indicates a negative correlation; a value of 1 indicates high similarity and 0 indicates no similarity. The formula for calculating similarity between ratings of two users X and Y , where \bar{X} and \bar{Y} denote the respective average of their ratings, is shown below:

$$\text{Pearson Correlation between user } X \text{ and user } Y = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2} \sqrt{\sum_{i=1}^n (Y_i - \bar{Y})^2}} \quad (2)$$

A common drawback of the collaborative filtering system is an active user could have high correlation with other users based on very few commonly rated items between them. The computed correlation between two users does not accurately represent their true correlation if we only have a few data points to compare their item ratings. As suggested in [1], the accuracy of the predictions could be improved if we can apply a significance weighting measure in the following way: if the number of commonly rated items (n) between the two users is less than a weight w then we can apply a significance weight of n/w . If the number of commonly rated items (n) between two users is more than the weight, then we apply a weight of 1 . The significance weights help us devalue the similarity weights that were based on fewer commonly-rated items as described in [1]. Ultimately, this enables us to produce a more reliable correlation

between two users. For our experiment, we have applied the following significance weight variations to each of our user-based and item-based processes for each similarity measure (euclidian distance and pearson):

- (1) significance weight of 1
- (2) significance weight of $n/25$ where n is the commonly rated items
- (3) significance weight of $n/50$ where n is the commonly rated items

In addition to applying the significance weights, it is crucial to choose the specific users whose ratings can be used in computing the predicted rating of the active user. As explained in [1], it is more accurate to use a subset of other similar users instead of the entire database to generate a prediction for the active user. For large-scale collaborative filtering recommender systems, it is time-consuming to use the entire user database to generate predictions. This is why it is important to selectively choose the best set of users to compute a prediction for the active user. As explained in [1], we have implemented the absolute threshold technique developed by *Shardanand and Maes et al., 1995* in our experiment. According to this method, we create our subset of other users by selecting only those users whose absolute correlation with the active user is greater than our given threshold. For our experiment, we have applied the following threshold values to each of our user-based and item-based processes for each similarity measure (euclidian distance and pearson):

- (1) threshold = 0
- (2) threshold = 0.3
- (3) threshold = 0.5

After choosing the best subset of other users, their ratings are utilised to compute a prediction for the active user. As mentioned earlier, the goal of our experiment is to identify the best set of recommender system algorithm. We would define the “best” recommendation system as the one that produces recommendations with the maximum coverage and the highest accuracy. The metrics used in our experiment to evaluate the quality of each recommender system algorithm can be explained in the following way:

- (1) Coverage: As defined in [1], coverage denotes the percentage of items for which the recommender system can compute predictions. In our experiment, coverage is shown as the number of items in place of a percentage value. It is desirable to have the maximum coverage as it implies that the prediction provided was computed based on almost all the item ratings existing in the dataset.
- (2) Accuracy: In our experiment, The accuracy of the predictions provided by a recommender system algorithm is evaluated using certain error metrics. For all the error metrics used in our experiment, the lowest error indicates the highest accuracy. The formula for each error metric is shown below:
 - (a) Mean Squared Error (MSE): The formula for calculating MSE where n denotes the number of ratings, Y_i denotes the actual rating and \bar{Y}_i denotes the predicted rating is the following:

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \bar{Y}_i)^2 \quad (3)$$

- (b) Root Mean Squared Error (RMSE): The formula for calculating RMSE where n denotes the number of ratings, R_i denotes the actual rating and P_i denotes the predicted rating is the following:

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (P_i - R_i)^2} \quad (4)$$

- (c) Mean Absolute Error (MAE): The formula for calculating MAE where n denotes the number of ratings, R_i denotes the actual rating and P_i denotes the predicted rating is the following:

$$MAE = \frac{1}{n} \sum_{i=1}^n |P_i - R_i| \quad (5)$$

Implementing the mathematical concepts and algorithms described above, we hypothesize that the item-based collaborative filtering process using the pearson similarity measure should produce the best recommendations with the maximum coverage and highest accuracy.

4 EXPERIMENTAL DESIGN

The objective for this experiment was to understand the effectiveness of different variations amongst item-based and user-based recommender systems. By applying different similarity weighting and similarity thresholds to each recommender system, we can produce different variations of each system. Then, we can further our analysis by comparing the changes in accuracy and coverage of each variation. Therefore, by implementing different variations and analyzing their accuracy and coverage, we can identify the most effective recommender system.

4.1 Experimental Design/Variations Requirements (36 variations total):

Below are the different parameters that we are running our tests with.

- Recommender System Algorithms: User-based and Item-based
- Similarity method: Distance, Pearson (+optional=Spearman or Jaccard or Tanimoto, etc.)
- Similarity significance weighting: None, $n/25$, $n/50$ (+optional= $n/75$, $n/100$, or variance weighting)
- Similarity threshold: >0 , >0.3 , >0.5 (+optional= >0.1 , >0.2 , >0.4)
- Rating prediction normalization: weighted (+optional=deviation-from-mean weighted or z-score predictions)
- Datasets: ML-100K (run some variations with critics first for testing; must provide Stats report for ML-100K)
- Evaluation: LOOCV (+optional=k-fold CV)
- Metrics: MSE, RMSE, MAE (report all three together)
- Tests of Hypothesis: (+optional=p-values for pairs of means or ANOVA for multiple means)

4.1.1 Datasets. For our experiment, we used the ML100k data from MovieLens. There are a total 943 users, 1664 movies, and 99693 ratings in this dataset. Ratings are provided in the range of 1 to 5. We then organized these data into a nested dictionary where each row is a user and each column is the movie that particular user has rated to conduct our study for user-based. In addition, we transformed the user-based dictionary to conduct our item-based study so that retrieving the rating data is faster. After setting up the rating dictionary, we created similarity matrices that are based on the parameters that were passed in: user-based/item-based, similarity weighting, and similarity threshold. The similarity data could then be used to predict items that the active user has yet to rate. To evaluate the accuracy and coverage of these predictions, we removed the active user's rating one by one and make a prediction based on the similarity matrix. Accuracy could then be measured by taking the error difference between the actual rating and our prediction rating using MSE, MAE and RMSE. A lower error metric would mean higher accuracy. We repeated this process for different parameters to analyze how different parameters impact the accuracy and coverage.

4.1.2 Layout of dataset statistics. Descriptive analytics data/Chart for ml-100K:

Number of users: 943

Number of items: 1664

Number of ratings: 99693

Overall average rating: 3.53 out of 5, and std dev of 1.13

Average item rating: 3.08 out of 5, and std dev of 0.78

Average user rating: 3.59 out of 5, and std dev of 0.44

User-item Matrix Sparsity: 93.65%

Average number of ratings per users: 105.718982, and std dev of 100.567291

Max number of ratings per users: 736

Min number of ratings per users: 19

Median number of ratings per users: 64.000000

Popular items – most rated		
Title	No. of Ratings	Average Rating
Star Wars (1977)	583	4.36
Contact (1997)	509	3.80
Fargo (1996)	508	4.16
Return of the Jedi (1983)	507	4.01
Liar Liar (1997)	485	3.16
Popular items – highest rated		
Title	Average Rating	No. of Ratings
Entertaining Angels: The Dorothy Day Story(1996)	5.00	1
Someone Else's America (1995)	5.00	1
Aiqing wansui (1994)	5.00	1
Star Kid (1997)	5.00	3
Marlene Dietrich: Shadow and Light (1996)	5.00	1

Overall best rated items (number of ratings ≥ 20)		
Title	Average Rating	No. of Ratings
Close Shave, A (1995)	4.49	112
Schindler's List (1993)	4.47	298
Wrong Trousers, The (1993)	4.47	118
Casablanca (1942)	4.46	243
Wallace & Gromit: The Best of Aardm	4.45	67
Shawshank Redemption, The (1994)	4.45	283
Rear Window (1954)	4.39	209
Usual Suspects, The (1995)	4.39	267
Star Wars (1977)	4.36	583
12 Angry Men (1957)	4.34	125
Third Man, The (1949)	4.33	72
Citizen Kane (1941)	4.29	198
Some Folks Call It a Sling Blade (1..	4.29	41
To Kill a Mockingbird (1962)	4.29	219
One Flew Over the Cuckoo's Nest (19.	4.29	264
Silence of the Lambs, The (1991)	4.29	390
North by Northwest (1959)	4.28	179
Godfather, The (1972)	4.28	179
Secrets & Lies (1996)	4.27	162
Good Will Hunting (1997)	4.26	198

4.1.3 Test Cases. For our experiment, we conducted the following tests for the implementation of the 36 variations.

- (1) User-based, Distance, similarity weight = None, similarity threshold = 0
- (2) User-based, Distance, similarity weight = $n/25$, similarity threshold = 0
- (3) User-based, Distance, similarity weight = $n/50$, similarity threshold = 0
- (4) User-based, Distance, similarity weight = None, similarity threshold = 0.3
- (5) User-based, Distance, similarity weight = $n/25$, similarity threshold = 0.3
- (6) User-based, Distance, similarity weight = $n/50$, similarity threshold = 0.3
- (7) User-based, Distance, similarity weight = None, similarity threshold = 0.5
- (8) User-based, Distance, similarity weight = $n/25$, similarity threshold = 0.5
- (9) User-based, Distance, similarity weight = $n/50$, similarity threshold = 0.5
- (10) Item-based, Distance, similarity weight = None, similarity threshold = 0
- (11) Item-based, Distance, similarity weight = $n/25$, similarity threshold = 0
- (12) Item-based, Distance, similarity weight = $n/50$, similarity threshold = 0
- (13) Item-based, Distance, similarity weight = None, similarity threshold = 0.3
- (14) Item-based, Distance, similarity weight = $n/25$, similarity threshold = 0.3
- (15) Item-based, Distance, similarity weight = $n/50$, similarity threshold = 0.3
- (16) Item-based, Distance, similarity weight = None, similarity threshold = 0.5
- (17) Item-based, Distance, similarity weight = $n/25$, similarity threshold = 0.5

- (18) Item-based, Distance, similarity weight = $n/50$, similarity threshold = 0.5
- (19) User-based, Pearson, similarity weight = None, similarity threshold = 0
- (20) User-based, Pearson, similarity weight = $n/25$, similarity threshold = 0
- (21) User-based, Pearson, similarity weight = $n/50$, similarity threshold = 0
- (22) User-based, Pearson, similarity weight = None, similarity threshold = 0.3
- (23) User-based, Pearson, similarity weight = $n/25$, similarity threshold = 0.3
- (24) User-based, Pearson, similarity weight = $n/50$, similarity threshold = 0.3
- (25) User-based, Pearson, similarity weight = None, similarity threshold = 0.5
- (26) User-based, Pearson, similarity weight = $n/25$, similarity threshold = 0.5
- (27) User-based, Pearson, similarity weight = $n/50$, similarity threshold = 0.5
- (28) Item-based, Pearson, similarity weight = None, similarity threshold = 0
- (29) Item-based, Pearson, similarity weight = $n/25$, similarity threshold = 0
- (30) Item-based, Pearson, similarity weight = $n/50$, similarity threshold = 0
- (31) Item-based, Pearson, similarity weight = None, similarity threshold = 0.3
- (32) Item-based, Pearson, similarity weight = $n/25$, similarity threshold = 0.3
- (33) Item-based, Pearson, similarity weight = $n/50$, similarity threshold = 0.3
- (34) Item-based, Pearson, similarity weight = None, similarity threshold = 0.5
- (35) Item-based, Pearson, similarity weight = $n/25$, similarity threshold = 0.5
- (36) Item-based, Pearson, similarity weight = $n/50$, similarity threshold = 0.5

4.1.4 Evaluation Metrics. We used MSE, MAE, and RMSE as our main tools to gauge the accuracy of the rating predictions. In terms of coverage, we ran all of our test cases to see the number of users that received recommendations from each algorithm.

5 RESULTS

- (1) **Recommendation Algorithm: User-based Collaborative Filtering, Item-based Collaborative Filtering:** Based on our results, we can observe that user-based collaborative filtering systems gives a higher coverage whereas item-based collaborative filtering systems provides higher accuracy with lower error metric values. (See Figure 4-6).
- (2) **Similarity method: Euclidean distance for RS, Pearson Correlation:** If we generally compare the two similarity measures, euclidian distance provides higher accuracy with lower error metric values whereas person correlation provides higher coverage (see Figure 9 and 10).
- (3) **Similarity significance weighting (None, $n/25$, $n/50$):** Generally speaking, a higher coverage was attained when a significance weight of 25 was applied to all the processes. (See Figure 7-10).
- (4) **Similarity threshold: > 0.0 , >0.3 , >0.5 :** A higher threshold of 0.3 and 0.5 returned lower error metric values for both user-distance and item-distance processes. (See Figure 1-6).
- (5) **Evaluation metric (MSE, RMSE, MAE):** All the error metrics returned no value for user-distance process (when the significance weight was 50) and for item-distance process (when the threshold was 0.5 for all significance weights).

5.1 Results Charts




Fig. 1. User-based Recommendation, y-axis: MSE, x-axis: simWtg (none, n/25, n/50)

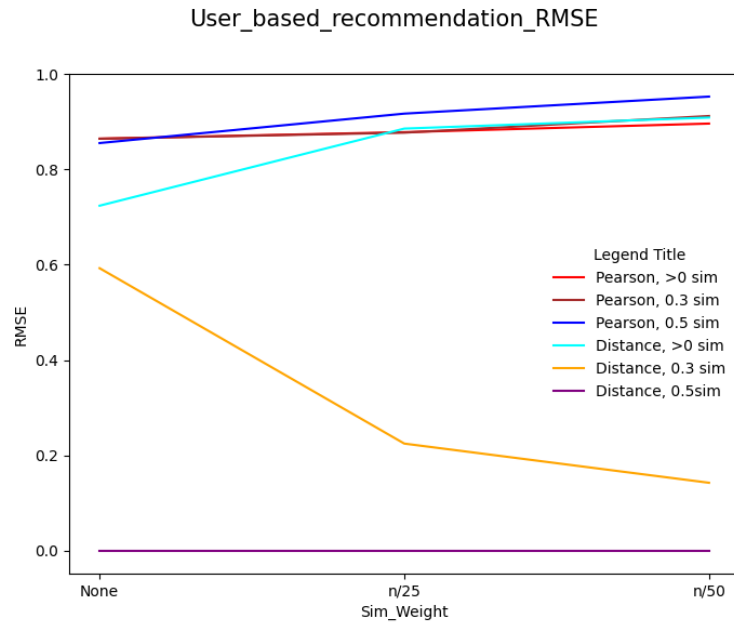


Fig. 2. User-based Recommendation, y-axis: RMSE, x-axis: simWtg (none, n/25, n/50)

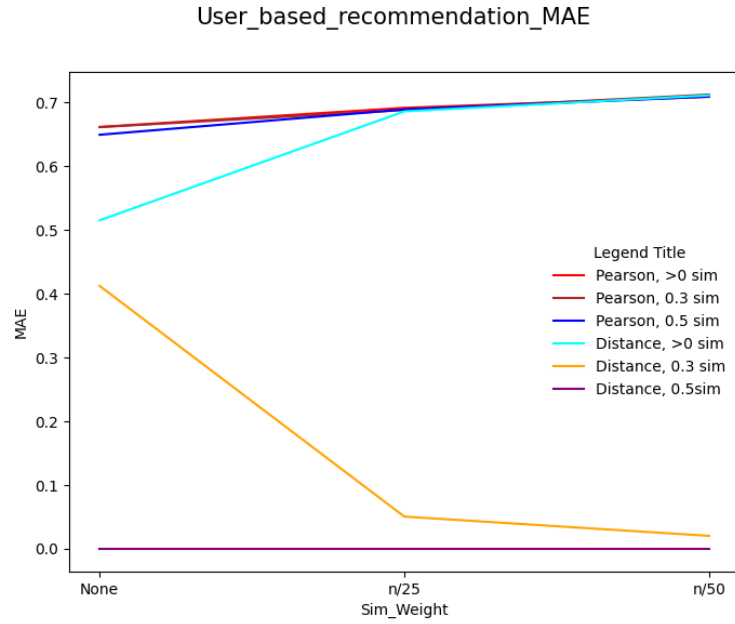


Fig. 3. User-based Recommendation, y-axis: MAE, x-axis: simWtg (none, n/25, n/50)

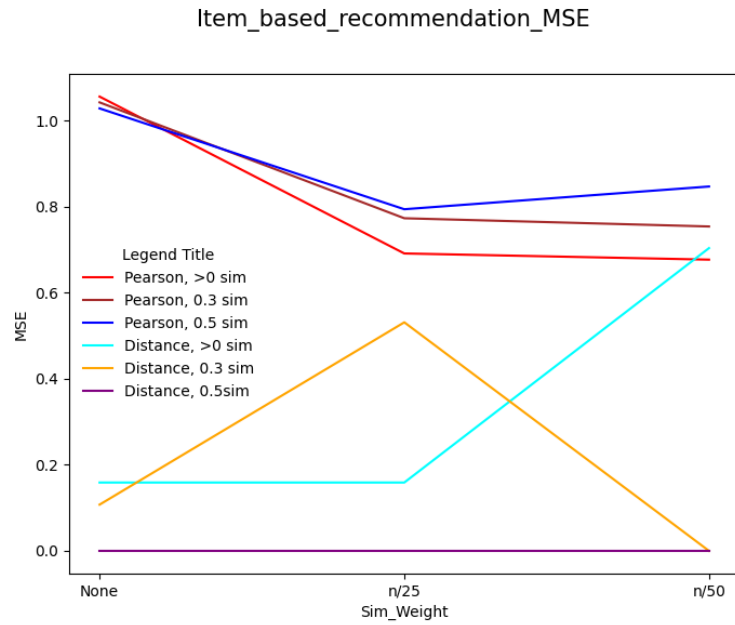


Fig. 4. Item-based Recommendation, y-axis: MSE, x-axis: simWtg (none, n/25, n/50)

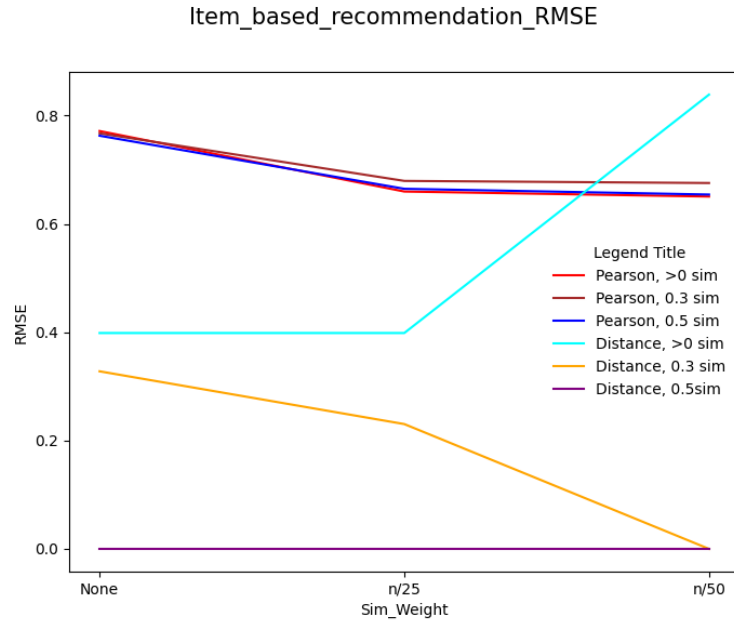


Fig. 5. Item-based Recommendation, y-axis: RMSE, x-axis: simWtg (none, n/25, n/50)




Fig. 6. Item-based Recommendation, y-axis: MAE, x-axis: simWtg (none, n/25, n/50)

Coverage vs. User-Distance (sim_wgt, threshold)

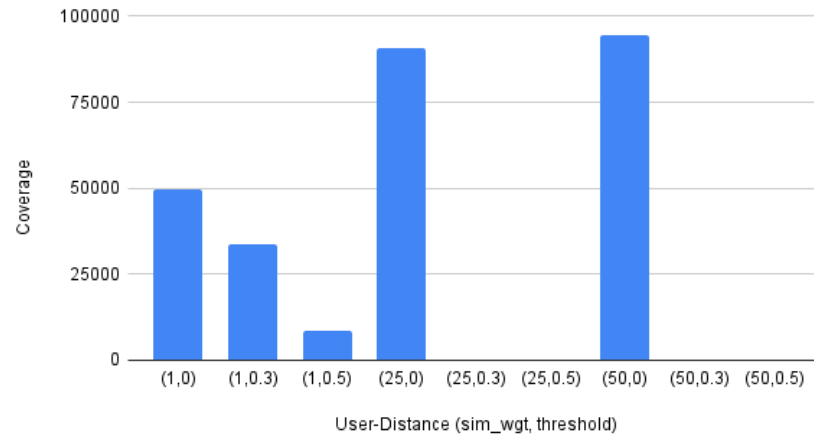


Fig. 7. Coverage vs. User-Distance

Coverage vs. User-Pearson (sim_wgt, threshold)

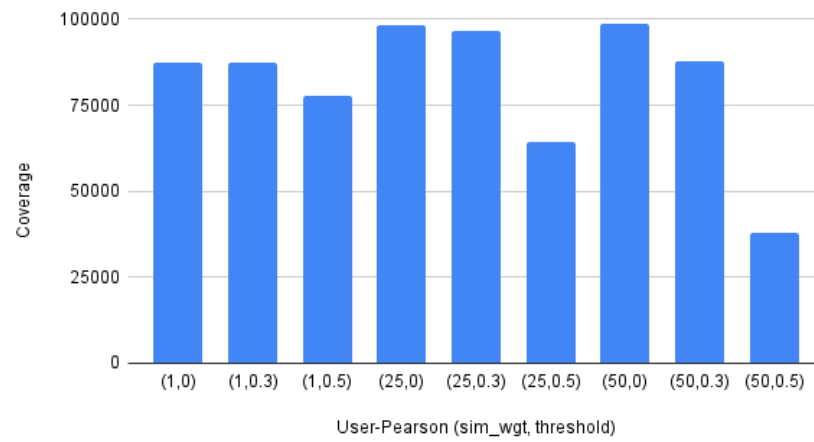


Fig. 8. Coverage vs. User-Pearson

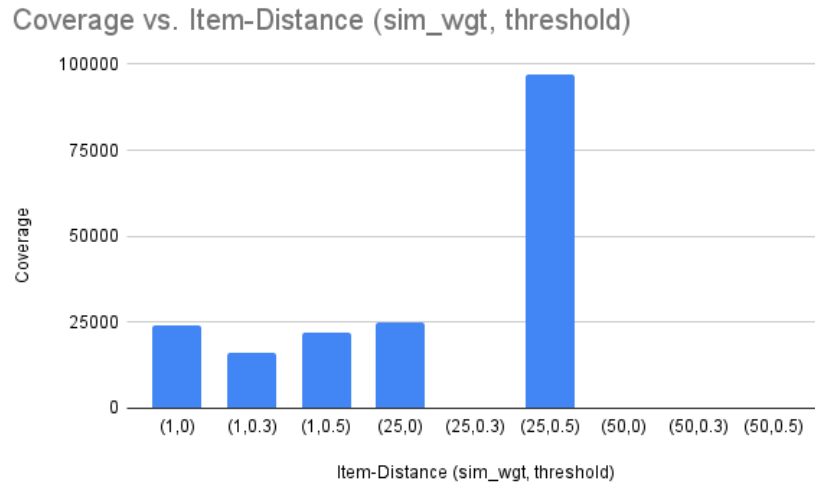


Fig. 9. Coverage vs. Item-Distance




Fig. 10. Coverage vs. Item-Pearson

6 DISCUSSION

Based on our results, our original hypotheses were partially rejected. Pearson was able to provide a better coverage for both user and item-based algorithms. However, contrary to what we have thought, larger similarity and larger threshold did not produce the best result. For example, as we can see with the results of when similarity weight = 50 and similarity threshold = 0.5, it produced very little coverage. By our definition of a good recommender system, we would want high coverage. From this, we can conclude that higher similarity weighting and threshold should be relevant to the size of the dataset. In our case, applying a high similarity weighting and threshold made our similar users and items decrease tremendously, to the point where no accurate prediction can be generated due to the lack of similarity. Additionally, we conducted a t-test by choosing a case where similarity weight was 25 and threshold was 0. This test helped us determine if there is a statistically significant difference between the means of our selected groups of data. We calculated a t-statistic and a p-value which allows us to compare MSE results (whether the means are equal) for three of our critics test cases. The results obtained are shown in the table below:

Critics	MSE	Coverage	P-Value
User-based CF, distance sim Vs User-based CF, pearson sim	0.78433 , 0.77109	90645, 98124	0.01974
Item-based CF, distance sim Vs Item-based CF, Pearson sim	0.81533 , 0.69092	97504, 85161	2.72804e-108
User-based CF, pearson sim Vs Item-based CF, Pearson sim	0.77109, 0.69092	98124 , 85161	5.53779e-57

Fig. 11. T-test results

After comparing the cases seen in figure 11, all of the comparisons reject the null hypothesis which states that the means are equal because the p-values were all lower than our chosen alpha value of 0.05.

Pearson correlation provides superior coverage and lower accuracy than Euclidean distance, while Euclidean distance provides higher accuracy at the expense of coverage. Because we consider higher coverage to be a more important metric, Pearson in that case fulfills the requirement better.

7 CONCLUSION

Our case study demonstrates that the user-based recommender system implemented with Pearson correlation, similarity weighting of $n/25$, and a threshold of 0 produces the best result overall. However, despite the general trend, there are some exceptions where the test cases of user-based Pearson and item-based Pearson did not produce conclusive results since there are minuscule differences in terms of accuracy and coverage between the two methods. Therefore, in our future research, we would like to conduct tests with different ranges of datasets to conclusively determine

which recommender algorithm performs better. For example, in our case where we defined our best recommender system as the one that provides the best coverage, which follows that the lower the threshold, the better the coverage. In this case, would our definition of a good recommender system offer the "best" recommender system? For further analysis, we can identify a threshold of coverage, so that once it reaches a good point of coverage, accuracy plays a more important weight on identifying the best recommender system. As we tested our datasets on different threshold values, coverage decreases dramatically as the threshold value increased. However, as our algorithms are applied to much larger datasets, it would not be practical to compute predictions using the ratings of all users. Therefore, the bigger concern lies between recognising what is more important for an effective recommender system: high coverage or higher accuracy. To strengthen our research, we plan to perform a more detailed analysis on a variety of thresholds in the future

8 ACKNOWLEDGMENTS

We are grateful to our CSC 381 professor Dr. Carlos Seminario who personally worked with us to help us navigate through the challenges in the project.

REFERENCES

- [1] Jonathan L. Herlocker, Joseph A. Konstan, Al Borchers, John Riedl, *An Algorithmic Framework for Performing Collaborative Filtering*, 1999
- [2] Badrul Sarwar, George Karypis, Joseph Konstan, and John Riedl, *Item-Based Collaborative Filtering Recommendation Algorithms*, 2001
- [3] Carlos Seminario, *Moodle Slides*, 2022